

Name: *Omar Rashad Salem*

Course: *CV - prof.Heba*

Assignment No.: 7

QUESTIONS

1) Explain the concept of a "dead ReLU" and why it can occur during training?

ans:

- Dead ReLU occurs when a ReLU neuron gets stuck outputting 0, causing it to become useless and potentially harm network performance.
- WHY?
 - drastic updates to weights specially during initialization can cause a dead ReLU
 - majority of input is negative combined with high learning rates also can cause a dead ReLU

2) What is dropout in neural networks. How does it prevent overfitting?

ans:

- is a regularization technique used in neural networks to prevent overfitting
- prevents overfitting by forcing the network to learn with fewer neurons, the dropout layer helps to prevent the network from memorizing the training data and instead learn more generalizable features. (learn don't memorize!)
- in Loss function:

$$L = \dots + \lambda R(W)$$

3) Given a neural network with a dropout rate of 0.3, calculate the probability that a specific neuron is dropped during a forward pass?

ans:

30%

4) Discuss the concept of early stopping in the context of training neural networks. What are the potential advantages and drawbacks?

ans: Early stopping aims to stop training once the model starts to overfit, preserving its generalizability.

- Advantages?
 - Reduces overfitting: learn the data don't memorize it!
 - computational efficiency: saves time!
 - Prevents model degradation: prevent model from becoming worse if continued training!
 - Disadvantages?
 - NONE: just kiddin! :D
 - Difficulty in debugging and interpreting
 - Difficulty in choosing stopping criteria
-

5) Explain the idea behind transfer learning and how it can be beneficial when training neural networks?

ans:

- Transfer learning: Taking a shortcut by using weights learned from another trained model on a specific dataset/s.
 - It's obviously very beneficial because it saves alot of time and boosts performance by widen the dataset we have with others datasets
-

Name: Omar Rashad Salem

Course: CV - prof.Heba

Assignment No.: 6

QUESTIONS

1) Explain the purpose of Batch Normalization in deep neural networks?

ans: Batch norm. layers help improving model overall performance by specifically boosting training stability and accelerating training process.

2) Describe the two main steps in the Batch Normalization process: normalization and scale/shift?

ans:

- 1. normalization: trying to get each batch `mean` to `0` and `variance` to `1` by subtracting the mean from the batch and deviding by `STD`

$$\mu_j = \frac{1}{N} \sum_{i=1}^N x_{i,j}$$
$$\sigma_j^2 = \frac{1}{N} \sum_{i=1}^N (x_{i,j} - \mu_j)^2$$
$$\therefore \hat{x}_{i,j} = \frac{x_{i,j} - \mu_j}{\sqrt{\sigma_j^2 + \epsilon}} (N \times D)$$

- 2. scale\shift: introducing new 2 learnable hyperparameters `gamma` and `beta` that finds the optimal scale and shift value of the normalized batch which makes the model converge faster

$$y = \gamma \hat{x} + \beta$$

3) How do the learnable parameters, gamma and beta, contribute to Batch Normalization?

ans: Gamma γ and beta β are parameters that are learned during training through backpropagation. These parameters enable the model to decide the optimal scale and shift for the normalized data.

$$y_{i,j} = \gamma_j \hat{x}_{i,j} + \beta_j$$

4) What was the groundbreaking contribution of AlexNet to the field of Convolutional Neural Networks?

ans:

- Introduced one of the first useful Deep not shallow ConvNet models (Achieved best scores in 2012 ImageNet Classification Challenge error = 16.4%)
 - Was one of the first models to use ReLU nonlinearities
 - Introduced very early normalization method called local response normalization
-

5) How does the VGG architecture differ from other ConvNet architectures in terms of filter size?

ans: Introduced a fixed rules to use:

- only 3x3 s=1 p=1 conv filters making scaling the model much easier
 - only max pooling layers 2x2 s=2 filters
-

6) Explain the concept of the inception module in GoogLeNet and its advantages?

ans:

- Inception module is a repeated local structure in GoogLeNet arch. works by simultaneously calculating multiple/all conv kernel sizes in same level then concatenate them effectively eliminating the need for kernel/filter size as a hyperparameter.
 - They also used bottleneck layer (filters) to reduce feature map dimensions before any expensive computation layers.
-

7) What is the main innovation introduced by ResNet to address training challenges in deep networks?

ans:

- Before ResNets after specific deep/depth ConvNets starts performing worse!
 - ResNets came to solve this issue by making learning the identity functions with high depth models easy and now deeper models again behaves as expected (better than shallower models)
 - After ResNet now we have two main blocks in Convnets a Plain block and a Residual block
-

Programming assignment:

```
In [ ]: # Omar rashad note: un-comment next line if first time using keras  
#! pip install keras
```

```
from sklearn.utils.multiclass import unique_labels  
from sklearn.model_selection import train_test_split  
from keras import Sequential  
from keras.applications import VGG16, ResNet50  
from keras.optimizers import SGD  
from keras.layers import Flatten,Dense #MAY NEED: BatchNormalization,Activation,Dropout  
from keras.utils import to_categorical  
  
#get the cifar10  
from keras.datasets import cifar10  
(x_train,y_train),(x_test,y_test)=cifar10.load_data()  
  
#get validation subset  
x_train,x_val,y_train,y_val=train_test_split(x_train,y_train,test_size=.3)  
  
#turn to onehot  
y_train=to_categorical(y_train)  
y_val=to_categorical(y_val)  
y_test=to_categorical(y_test)  
  
#Defining the VGG and ResNet50  
vgg16_base_model = VGG16(include_top=False,weights='imagenet',input_shape=(32,32,3),classes=y_train.shape[1])  
resnet50_base_model = ResNet50(include_top=False,weights='imagenet',input_shape=(32,32,3),classes=y_train.shape[1])  
  
#create our own modified model ( append our dens layer at end)  
model= Sequential()  
model.add(vgg16_base_model)  
model.add(Flatten())  
  
model2= Sequential()  
model2.add(resnet50_base_model)  
model2.add(Flatten())  
  
#after creating ? yes add the layers!  
model.add(Dense(1024,activation='relu'),input_dim=512))  
model.add(Dense(512,activation='relu')))  
model.add(Dense(256,activation='relu')))  
model.add(Dense(128,activation='relu')))  
model.add(Dense(10,activation='softmax')) #This is the classification Layer  
  
model2.add(Dense(1024,activation='relu'),input_dim=512))  
model2.add(Dense(512,activation='relu')))
```

```
model2.add(Dense(256,activation='relu')))
model2.add(Dense(128,activation='relu')))
model2.add(Dense(10,activation='softmax')) #This is the classification Layer

#Let's see it now! ORS#
print(f'modified VGG model Summary: \n')
model.summary()
print(f'modified ResNet model Summary: \n')
model2.summary()

#hype hype :D hyperparameters
batch_size= 100
epochs=5
learn_rate=.001
sgd=SGD(learning_rate=learn_rate,momentum=.9,nesterov=False)

#all set! COMPILE!
model.compile(optimizer=sgd,loss='categorical_crossentropy',metrics=['accuracy'])
model2.compile(optimizer=sgd,loss='categorical_crossentropy',metrics=['accuracy'])

#train and save train history
print(f'\n\n\n START TRAINING OUR ResNet \n\n\n')
history2 = model2.fit(x_train, y_train, batch_size= batch_size, epochs= epochs, validation_data=(x_val, y_val))
print(f'\n\n\n START TRAINING OUR VGG \n\n\n')
history = model.fit(x_train, y_train, batch_size= batch_size, epochs= epochs, validation_data=(x_val, y_val))
```

modified VGG model Summary:

Model: "sequential_7"

Layer (type)	Output Shape	Param #
=====		
vgg16 (Functional)	(None, 1, 1, 512)	14714688
flatten_7 (Flatten)	(None, 512)	0
dense_35 (Dense)	(None, 1024)	525312
dense_36 (Dense)	(None, 512)	524800
dense_37 (Dense)	(None, 256)	131328
dense_38 (Dense)	(None, 128)	32896
dense_39 (Dense)	(None, 10)	1290
=====		
Total params:	15,930,314	
Trainable params:	15,930,314	
Non-trainable params:	0	

modified ResNet model Summary:

Model: "sequential_8"

Layer (type)	Output Shape	Param #
=====		
resnet50 (Functional)	(None, 1, 1, 2048)	23587712
flatten_8 (Flatten)	(None, 2048)	0
dense_40 (Dense)	(None, 1024)	2098176
dense_41 (Dense)	(None, 512)	524800
dense_42 (Dense)	(None, 256)	131328
dense_43 (Dense)	(None, 128)	32896
dense_44 (Dense)	(None, 10)	1290
=====		
Total params:	26,376,202	
Trainable params:	26,323,082	

Non-trainable params: 53,120

In []:

```
#get the history
print(f'\n\nOur VGG model train history: ')
for info in history.history: print(f'Epochs {info} : {history.history[info]}')

print(f'\n\nOur ResNet model train history: ')
for info in history2.history: print(f'Epochs {info} : {history2.history[info]}')

print('\n\n')
print('Total Train Time for Both models = 88m 30.8s')
print('\n\n')

print('E V A L U A T E')
#evaluate
loss, accuracy = model.evaluate(x_test, y_test)
print(f'Our VGG model final results on cifar10 dataset: \n loss= {loss}\n accuracy= {accuracy}\n\n')
loss2, accuracy2 = model2.evaluate(x_test, y_test)
print(f'Our ResNet model final results on cifar10 dataset: \n loss= {loss2}\n accuracy= {accuracy2}')

#DONE!
```

Our VGG model train history:

```
Epochs loss : [2.331294536590576, 2.302603244781494, 2.302593946456909, 2.3025898933410645, 2.3025882244110107]
Epochs accuracy : [0.09759999811649323, 0.09868571162223816, 0.09757142513990402, 0.09991428256034851, 0.09797143191099167]
Epochs val_loss : [2.3026533126831055, 2.30268931388855, 2.3027286529541016, 2.302746295928955, 2.3027658462524414]
Epochs val_accuracy : [0.09880000352859497, 0.09753333032131195, 0.09753333032131195, 0.09746666997671127, 0.09746666997671127]
```

Our ResNet model train history:

```
Epochs loss : [1.592706561088562, 0.8918847441673279, 0.6395752429962158, 0.45397523045539856, 0.342894583940506]
Epochs accuracy : [0.43308570981025696, 0.689542829990387, 0.7785428762435913, 0.8407999873161316, 0.8809428811073303]
Epochs val_loss : [1.126675009727478, 0.8955151438713074, 0.856717050075531, 0.8438029289245605, 0.8935683369636536]
Epochs val_accuracy : [0.6024666428565979, 0.6928666830062866, 0.707733331108093, 0.73253333568573, 0.7376000285148621]
```

Total Train Time for Both models = 88m 30.8s

E V A L U A T E

```
313/313 [=====] - 18s 57ms/step - loss: 2.3026 - accuracy: 0.1000
```

Our VGG model final results on cifar10 dataset:

```
loss= 2.3026132583618164
accuracy= 0.10000000149011612
```

```
313/313 [=====] - 15s 48ms/step - loss: 0.8945 - accuracy: 0.7362
```

Our ResNet model final results on cifar10 dataset:

```
loss= 0.894540548324585
accuracy= 0.7361999750137329
```